Tadisetty Sai Yashwanth

Bengaluru, KA, IN | taddishetty34@gmail.com | +91 7795348927 | saiyashwanth.com | linkedin.com/in/saiyashwanth29

github.com/theyashwanthsai

Education

PES University, Btech in Computer Science and Engineering 09/2021 – 05/2025

- GPA: 8.43/10.0
- **Teaching Assistant - Intro to Generative AI**: Collaborated with faculty to enhance course materials, lectures, assignments, and assessments by integrating student feedback and aligning content with recent advancements in generative AI.

Experience

AI Engineer, Vuhosi – UK, Remote Aug 2024 – present

Leading AI initiatives to build the main product: copilot for investors and startup founders. Designed and deployed core agentic RAG systems using CrewAI. Spearheading the development of production-grade AI agents powering the product.

AI Dev - Contract Gig, Plotweaver Apr 2025 – June 2025

Contracted via Vuhosi to support rapid prototyping of the MVP product. Collaborated closely with the team to build AI-first features aligned with storytelling workflows.

AI Dev - Contract Gig, Carbonnote Dec 2024 – Feb 2025

Contracted via Vuhosi to help finalize and ship GreenIQ, an AI-driven sustainability analytics tool. Contributed to both product development and co-authored a research paper on its technical foundation.

Summer Intern, Composio – Bengaluru, In May 2024 – July 2024

Worked in DevRel to build and showcase AI products using Composio. Created 15+ guides, a custom OpenAI GPT, and multiple projects featured in documentation and social media, contributing to repo growth from 50 to 2K stars.

Dev Intern, VertexCover – Remote Dec 2023 – Jan 2024

Collaborated on GateGPT, an open-source third-party auth system for custom GPTs, integrating FastAPI, Pydantic, and Google OAuth to enhance security and UX.

Research and Book Publications

Book: Multi Agent Systems in Action - Co Authoring Ongoing

Under Manning Publications. Writing this book with the founder of vuhosi.

Paper: On the Structure of Floating-Point Noise in Batch-Invariant GPU Matrix Multiplication (Pre Print)

Recently finished the study, looking for A* Conferences to submit. Arxiv paper will be out soon.

Paper: Reasoning Under Pressure: LLMs in Competitive Pokémon Battles (Pre Print)

Accepted at **Neurips** Law 2025 Workshop.

Paper: Real Time Child Abduction And Detection System (Pre Print)

Accepted and presented at the 15th IEEE International Conference on Intelligent Systems: Theories and Applications (SITA 2025). Oct 20, 2025.

Paper: GreenIQ: A Deep Search Platform for Comprehensive Carbon Market Analysis and Automated Report Generation (Pre Print)

Accepted and Presented at the 10th International Conference on Machine Learning Technologies (ICMLT 2025). May 2025.

Paper: Advancing Cricket Narratives: AI-Enhanced Advanced Journaling in the IPL Using Language Models (link)

Accepted and Presented at the 10th IEEE CONECCT. July 14, 2024.

Projects

Devyan (270+ stars on Github) Python, Crewai. June 2024

Built a multi-agent system that automates software development from a single prompt, orchestrating a team of agents: Architect, Programmer, Tester, and Reviewer to collaboratively plan, implement, test, and refine software projects end-to-end. This was built before the likes of bolt or lovable existed, as a fun side project. (Github) (Tweet) (Video)

Pokemon Tournament Bench Python, React. Feb 2025

Developed a framework where LLMs battle each other in simulated Pokémon-style matches, with a React interface visualizing gameplay. Designed to study strategic reasoning and model behavior under game constraints. Accepted at neurips 2025 workshop. (Github) (Demo)

MREB: Multimodal Benchmark for SMLs Python, Ollama. May 2025

Created an open-source benchmark to evaluate small LLMs running locally via Ollama across 4 categories: logic, coding, ethics, and multimodal tasks. Includes 100+ test cases, automated scoring, visual dashboards, and a leaderboard for consistent model evaluation. Designed to promote transparent benchmarking for local AI models. Currently adding tools bench. (Github) (Blog) (Video)

Zenscript Python. Sept 2023

Designed and implemented a toy programming language in Python to understand compiler and interpreter design. Built core components including a lexer, parser, AST evaluator, and runtime environment, enabling execution of user-defined logic and control flow. I built this in my 5th sem, as a fun project but later I became serious about the subject. (Github)

OS-2 C, Assembly. July 2025

Vibe-coded a simple 32-bit operating system from scratch in C and Assembly purely for fun as I was bored with prompting and python. Implemented a custom bootloader that switches to protected mode and rewrote the kernel in C. Work-in-progress project featuring dev logs that document the process. Long-term goal is to get Doom running on the OS: 5 years from now. (Github) (Devlog)

Achievements & Notable Work

Reviewer at Neurips ER Workshop

Invited to serve as a Reviewer for the NeurIPS 2025 ER Workshop by the program committee, contributing to the peer review process for one of the premier AI research conferences, enhancing academic rigor and community impact.

Public Speaking

Delivered two technical talks at Microsoft Reactor on building and deploying AI agents, engaging a diverse developer audience.

Technical Writing & Knowledge Sharing

Authored multiple technical essays on AI agents on X, reaching 300K+ impressions and sparking insightful community discussions. View Essays
Conducted workshops for school students to educate them about Hackathons and AI. Link

Academic & Community Engagement

- Recieved MRD and DAC scholarships at PES University
- Won Best Capstone Project Award in Panel 6 at Capstone Project Mela by Department of CSE, PES University. Link
- Mentor at HackerSpace (HSP), nurturing top tech talent at PES University.
- Member of ACM PESU ECC; volunteered at multiple hackathons and technical events. Team Lead at AIEP 2022
- Core member of PES Animal Welfare Club, leading dog care and feeding initiatives on campus.

Hackathons

- Judge, OPQtech Summer AI Hackathon 2025
- Mentor, ACM Hacknight 2023
- 3rd Place, Kalpana – Hack for Humanity 2024. Link
- 3rd Place, MLH Global Hack Week May 2023. Link

Leadership

Founder, Turi Labs Started Nov 2024

Turilabs is a lean, indie AI research lab dedicated to building weird, wonderful, and sometimes outright crazy things. We're driven by raw curiosity to explore the next generation of intelligent systems that can learn, adapt, and evolve. [Our website]
Key initiatives include:

- Developed an in-house agentic framework powered by OpenAI's Assistants API to orchestrate multi-agent collaboration and streamline complex research workflows. Paper coming soon
- Built a quirky benchmark inspired by the Pokémon League, each LLM acts as a "trainer" competing for the title of *Pokémon Master*. Accepted at NeurIPS 2025 Workshop: LAW. (Preprint)
- Currently ideating and brainstorming on a compact world model trained entirely on Pokémon game's overworld to learn more about simulated worlds.
- Leading a diverse but closed research community spanning from first-year students to a 65-year-old reinforcement learning enthusiast, fostering an inclusive and curiosity-driven ecosystem.
- Around 5 research projects nearing publication, exploring themes around agentic systems and LLM behavior in simulated environments.