Tadisetty Sai Yashwanth

Bengaluru, KA, IN | taddishetty34@gmail.com | +91 7795348927 | saiyashwanth.com | Google Scholar
linkedin.com/in/saiyashwanth29 | github.com/theyashwanthsai

Research Interests

Reinforcement learning environments and post-training, evaluation of LLM agents through closed-loop, execution-grounded benchmarks, and the reliability of LLM reasoning, from logical inconsistency to GPU-level numerical non-determinism.

Education

PES University, B.Tech in Computer Science and Engineering Sep 2021 – May 2025

- GPA: 8.43/10.0. Received MRD and DAC merit scholarships.
- **Teaching Assistant - Intro to Generative AI**: Collaborated with faculty to enhance course materials, lectures, assignments, and assessments by integrating student feedback and aligning content with recent advancements in generative AI.
- **Best Capstone Project Award** in Panel 6 at Capstone Project Mela, Department of CSE. The project was published at IEEE SITA 2025.

Research Publications

Paper: RLRank: Distilling Offline Oracles into Online Policies for Document Reranking (Link)
S. Palchaudhuri, **T. S. Yashwanth**. Accepted at **ICML** 2026 DEMO Workshop (Poster).

Paper: Executable Ground Truth: A Closed-Loop Benchmark for Evaluating LLM Agents on Microservice Incident Remediation (Link)
C. Dhatri, **T. S. Yashwanth**. Accepted at **ICML** 2026 CTB Workshop (Poster).

Paper: Compositional Skill Acquisition in Agentic Pipelines via Reinforcement Learning and Knowledge Distillation (Link)
T. S. Yashwanth. Accepted at **ICML** 2026 CompLearn Workshop (Poster).

Paper: Benchmarking Logical Reasoning Inconsistencies in Local Large Language Models: Evidence from Multi-Domain Evaluation (Link)
T. S. Yashwanth, C. Dhatri. Accepted at **ICLR** 2026 Workshop on Logical Reasoning of Large Language Models.

Paper: On the Structure of Floating-Point Noise in Batch-Invariant GPU Matrix Multiplication (Link)
T. S. Yashwanth. Accepted at **ICLR** 2026 Sci4DL Workshop.

Paper: Reasoning Under Pressure: LLMs in Competitive Pokémon Battles (Link)
T. S. Yashwanth, C. Dhatri. Accepted at **NeurIPS** 2025 ER Workshop.

Paper: Real Time Child Abduction And Detection System (Link)
T. S. Yashwanth, Y. S. Royal, M. Kashyap, V. R. Shreya, D. K. N. Accepted and presented at the 15th IEEE International Conference on Intelligent Systems: Theories and Applications (SITA 2025). Oct 2025.

Paper: GreenIQ: A Deep Search Platform for Comprehensive Carbon Market Analysis and Automated Report Generation (Link)
O. Fagbohun, **T. S. Yashwanth**, A. S. Akintola, et al. Accepted and presented at the 10th International Conference on Machine Learning Technologies (ICMLT 2025). May 2025.

Paper: Advancing Cricket Narratives: AI-Enhanced Advanced Journaling in the IPL Using Language Models (Link)
S. Sarkar, **T. S. Yashwanth**, A. Giri. Accepted and presented at the 10th IEEE CONECCT. July 2024.

Book: Agent Engineering: Building, Evaluating, and Operating Reliable AI Agents Under Contract
Co-author with O. Fagbohun, under contract with Apress, covering the engineering principles and practices required to build, evaluate, and operate reliable AI agents.

Experience

AI Engineer, StartupFuel – Canada, Remote Dec 2025 – Jun 2026
Built AI agents for finance and investments, including a newsletter agent, a market research agent, and a gap analysis agent, and post-trained language models for the same domain. Built a product monitoring layer on PostHog and an AI observability layer for tracing agent behavior in production, and wrote evals for each agent to measure output quality and catch regressions.

AI Engineer, Vuhosi – UK, Remote Aug 2024 – Nov 2025
Led AI initiatives for the main product, a copilot for investors and startup founders. Designed and deployed its core agentic RAG systems and production-grade AI agents, including an automated pitch-deck evaluator: founders email their deck and, within 60 minutes, receive a report that scores the startup on 27 criteria, covering strengths, weaknesses, risks, an investment-readiness

decision, and investors matched to their niche. (Demo) Through Vuhosi, also took on client contracts: helped ship GreenIQ for Carbonnote (Dec 2024 – Feb 2025), an AI-driven carbon-market analytics tool, and co-authored its ICMLT 2025 paper. Also prototyped AI-first storytelling features for Plotweaver's MVP (Apr 2025 – Jun 2025).

Summer Intern, Composio – Bengaluru, IN May 2024 – Jul 2024
Worked in DevRel to build and showcase AI products using Composio. Created 15+ guides, a custom OpenAI GPT, and multiple projects featured in documentation and social media, contributing to repo growth from 50 to 2K stars.

Leadership

Founder and Research Lead, Turi Labs Nov 2024 – Present
Turi Labs is a lean, indie AI research lab working on reinforcement learning, agentic systems, evaluation, and LLM reasoning.

- Led research that produced **6 workshop papers at NeurIPS, ICLR, and ICML**, spanning AI agents, LLM reasoning, reinforcement learning, benchmarks, and GPU computation.
- Built a research community of **25+ members** and mentored its part-time researchers, from first-year undergraduates to industry veterans. Scoped projects, guided experiments, and co-wrote papers.
- Secured a **$5,000 research grant from Thinking Machines Lab**, **$10,000 in Azure credits** through Microsoft for Startups, and Claude Code and API credits through Claude for Startups.
- Conducted **2 offline meetups** bringing together members of the AI research and engineering community.

Projects

Harness Dojo 2026 – Present

A LeetCode and CodeCrafters-style platform for **harness engineering**, the discipline of building reliable scaffolding, evals, and environments around AI models and agents. Launched the public beta with **15 active users** and the first **5 problems** live. (Website)

jevals and evalkit Python. 2026

Open-source tools for evaluating LLMs and agents. jevals grades outputs with a calibrated decision model, so every result comes with a confidence score and only uncertain cases go to a human or LLM judge. evalkit makes evals a first-class part of the development loop by keeping offline and online evals, traces, and agent versions in git. (jevals) (evalkit)

Devyan (290+ stars on GitHub) Python, CrewAI. June 2024

Built a multi-agent system that automates software development from a single prompt, orchestrating a team of agents (Architect, Programmer, Tester, and Reviewer) to plan, implement, test, and refine software projects end to end. (Github) (Tweet) (Video)

Grants, Awards & Notable Work

Grants and Partnerships

- $5,000 research grant from Thinking Machines Lab for independent research at Turi Labs.
- Microsoft for Startups partner ($10,000 in Azure credits) and Claude for Startups partner (Claude Code and API credits).

Reviewer at NeurIPS ER Workshop

Invited by the program committee to serve as a reviewer for the NeurIPS 2025 ER Workshop.

Public Speaking

Delivered two technical talks at Microsoft Reactor on building and deploying AI agents. (Talk 1) (Talk 2)

Technical Writing & Knowledge Sharing

Authored multiple technical essays on AI agents on X, reaching 300K+ impressions and sparking insightful community discussions. (View Essays)

Mentoring

Mentor at HackerSpace (HSP), nurturing top tech talent at PES University.

Hackathons

- Judge, TechFusion 2.0 hosted by the Department of CSE at Dayananda Sagar Academy of Technology and Management, Bengaluru.
- Judge, OPQtech Summer AI Hackathon 2025
- Mentor, ACM Hacknight 2023
- 3rd Place, Kalpana – Hack for Humanity 2024. Link
- 3rd Place, MLH Global Hack Week May 2023. Link